

Question #51

Topic 1

A Snowflake Architect is designing a multi-tenant application strategy for an organization in the Snowflake Data Cloud and is considering using an Account Per Tenant strategy.

Which requirements will be addressed with this approach? (Choose two.)

- A. There needs to be fewer objects per tenant.
- B. Security and Role-Based Access Control (RBAC) policies must be simple to configure. Most Voted
- C. Compute costs must be optimized.
- D. Tenant data shape may be unique per tenant. Most Voted
- E. Storage costs must be optimized.

**Correct Answer:** CE

Question #52

Topic 1

An Architect has been asked to clone schema STAGING as it looked one week ago, Tuesday June 1st at 8:00 AM, to recover some objects.

The STAGING schema has 50 days of retention.

The Architect runs the following statement:

```
CREATE SCHEMA STAGING_CLONE CLONE STAGING at (timestamp => '2021-06-01 08:00:00');
```

The Architect receives the following error: Time travel data is not available for schema STAGING. The requested time is either beyond the allowed time travel period or before the object creation time.

The Architect then checks the schema history and sees the following:

```
CREATED_ON|NAME|DROPPED_ON -  
2021-06-02 23:00:00 | STAGING | NULL  
2021-05-01 10:00:00 | STAGING | 2021-06-02 23:00:00
```

How can cloning the STAGING schema be achieved?

- A. Undrop the STAGING schema and then rerun the CLONE statement.
- B. Modify the statement: CREATE SCHEMA STAGING\_CLONE CLONE STAGING at (timestamp => '2021-05-01 10:00:00');
- C. Rename the STAGING schema and perform an UNDROP to retrieve the previous STAGING schema version, then run the CLONE statement. Most Voted
- D. Cloning cannot be accomplished because the STAGING schema version was not active during the proposed Time Travel time period.

**Correct Answer:** C

What are purposes for creating a storage integration? (Choose three.)

- A. Control access to Snowflake data using a master encryption key that is maintained in the cloud provider's key management service.
- B. Store a generated identity and access management (IAM) entity for an external cloud provider regardless of the cloud provider that hosts the Snowflake account. Most Voted
- C. Support multiple external stages using one single Snowflake object. Most Voted
- D. Avoid supplying credentials when creating a stage or when loading or unloading data. Most Voted
- E. Create private VPC endpoints that allow direct, secure connectivity between VPCs without traversing the public internet.
- F. Manage credentials from multiple cloud providers in one single Snowflake object.

**Correct Answer:** BCD

A healthcare company is deploying a Snowflake account that may include Personal Health Information (PHI). The company must ensure compliance with all relevant privacy standards.

Which best practice recommendations will meet data protection and compliance requirements? (Choose three.)

- A. Use, at minimum, the Business Critical edition of Snowflake. Most Voted
- B. Create Dynamic Data Masking policies and apply them to columns that contain PHI. Most Voted
- C. Use the Internal Tokenization feature to obfuscate sensitive data.
- D. Use the External Tokenization feature to obfuscate sensitive data. Most Voted
- E. Rewrite SQL queries to eliminate projections of PHI data based on current\_role().
- F. Avoid sharing data with partner organizations.

**Correct Answer:** ABD

There are two databases in an account, named fin\_db and hr\_db which contain payroll and employee data, respectively. Accountants and Analysts in the company require different permissions on the objects in these databases to perform their jobs. Accountants need read-write access to fin\_db but only require read-only access to hr\_db because the database is maintained by human resources personnel.

An Architect needs to create a read-only role for certain employees working in the human resources department.

Which permission sets must be granted to this role?

- A. USAGE on database hr\_db, USAGE on all schemas in database hr\_db, SELECT on all tables in database hr\_db Most Voted
- B. USAGE on database hr\_db, SELECT on all schemas in database hr\_db, SELECT on all tables in database hr\_db
- C. MODIFY on database hr\_db, USAGE on all schemas in database hr\_db, USAGE on all tables in database hr\_db
- D. USAGE on database hr\_db, USAGE on all schemas in database hr\_db, REFERENCES on all tables in database hr\_db

**Correct Answer:** A

An Architect runs the following SQL query:

```
SELECT  
    METADATA$FILENAME,  
    METADATA$FILE_ROW_NUMBER  
FROM @FILEROWS/Food_Reviews.csv  
    (file_format=CSV_N)
```

How can this query be interpreted?

- A. FILEROWS is a stage. FILE\_ROW\_NUMBER is line number in file. Most Voted
- B. FILEROWS is the table. FILE\_ROW\_NUMBER is the line number in the table.
- C. FILEROWS is a file. FILE\_ROW\_NUMBER is the file format location.
- D. FILERONS is the file format location. FILE\_ROW\_NUMBER is a stage.

**Correct Answer: A**

An Architect entered the following commands in sequence:

```
CREATE DATABASE SANDBOX;  
CREATE ROLE INTERN;  
CREATE TABLE SANDBOX.PUBLIC.AGENDA (ID INT, ITEMS STRING);  
GRANT SELECT ON ALL TABLES IN SCHEMA SANDBOX.PUBLIC TO ROLE INTERN;  
GRANT ROLE INTERN TO USER USER1;
```

USER1 cannot find the table.

Which of the following commands does the Architect need to run for USER1 to find the tables using the Principle of Least Privilege? (Choose two.)

- A. GRANT ROLE PUBLIC TO ROLE INTERN;
- B. GRANT USAGE ON DATABASE SANDBOX TO ROLE INTERN;
- C. GRANT USAGE ON SCHEMA SANDBOX.PUBLIC TO ROLE INTERN;
- D. GRANT OWNERSHIP ON DATABASE SANDBOX TO USER INTERN;
- E. GRANT ALL PRIVILEGES ON DATABASE SANDBOX TO ROLE INTERN;

**Correct Answer: BC**

A DevOps team has a requirement for recovery of staging tables used in a complex set of data pipelines. The staging tables are all located in the same staging schema. One of the requirements is to have online recovery of data on a rolling 7-day basis.

After setting up the DATA\_RETENTION\_TIME\_IN\_DAYS at the database level, certain tables remain unrecoverable past 1 day.

What would cause this to occur? (Choose two.)

- A. The staging schema has not been setup for MANAGED ACCESS.
- B. The DATA\_RETENTION\_TIME\_IN\_DAYS for the staging schema has been set to 1 day.
- C. The tables exceed the 1 TB limit for data recovery.
- D. The staging tables are of the TRANSIENT type.
- E. The DevOps role should be granted ALLOW\_RECOVERY privilege on the staging schema.

**Correct Answer:** BD

Consider the following COPY command which is loading data with CSV format into a Snowflake table from an internal stage through a data transformation query.

```
copy into home_sales(city, zip, sale_date, price)
from (select t.$1, t.$2, t.$6, t.$7 from @mystage/sales.csv.qz t)
file_format =
(
  format_name = mycsvformat
  empty_field_as_null = true
  field_optionally_enclosed_by = ''
)
validation_mode - return_all_errors
;
```

This command results in the following error:

SQL compilation error: invalid parameter 'validation\_mode'

Assuming the syntax is correct, what is the cause of this error?

- A. The VALIDATION\_MODE parameter supports COPY statements that load data from external stages only.
- B. The VALIDATION\_MODE parameter does not support COPY statements with CSV file formats.
- C. The VALIDATION\_MODE parameter does not support COPY statements that transform data during a load. Most Voted
- D. The value return\_all\_errors of the option VALIDATION\_MODE is causing a compilation error.

**Correct Answer:** C

A company is storing large numbers of small JSON files (ranging from 1-4 bytes) that are received from IoT devices and sent to a cloud provider. In any given hour, 100,000 files are added to the cloud provider. What is the MOST cost-effective way to bring this data into a Snowflake table?

- A. An external table
- B. A pipe**
- C. A stream
- D. A copy command at regular intervals

**Correct Answer:** B

A company has a Snowflake account named ACCOUNTA in AWS us-east-1 region. The company stores its marketing data in a Snowflake database named MARKET\_DB. One of the company's business partners has an account named PARTNERB in Azure East US 2 region. For marketing purposes the company has agreed to share the database MARKET\_DB with the partner account.

Which of the following steps MUST be performed for the account PARTNERB to consume data from the MARKET\_DB database?

- A. Create a new account (called AZABC123) in Azure East US 2 region. From account ACCOUNTA create a share of database MARKET\_DB, create a new database out of this share locally in AWS us-east-1 region, and replicate this new database to AZABC123 account. Then set up data sharing to the PARTNERB account.
- B. From account ACCOUNTA create a share of database MARKET\_DB, and create a new database out of this share locally in AWS us-east-1 region. Then make this database the provider and share it with the PARTNERB account.
- C. Create a new account (called AZABC123) in Azure East US 2 region. From account ACCOUNTA replicate the database MARKET\_DB to AZABC123 and from this account set up the data sharing to the PARTNERB account. Most Voted**
- D. Create a share of database MARKET\_DB, and create a new database out of this share locally in AWS us-east-1 region. Then replicate this database to the partner's account PARTNERB.

**Correct Answer:** C -

What does a Snowflake Architect need to consider when implementing a Snowflake Connector for Kafka?

- A. Every Kafka message is in JSON or Avro format. Most Voted
- B. The default retention time for Kafka topics is 14 days.
- C. The Kafka connector supports key pair authentication, OAUTH, and basic authentication (for example, username and password).
- D. The Kafka connector will create one table and one pipe to ingest data for each topic. If the connector cannot create the table or the pipe it will result in an exception.**

**Correct Answer:** D

A group of Data Analysts have been granted the role ANALYST\_ROLE. They need a Snowflake database where they can create and modify tables, views, and other objects to load with their own data. The Analysts should not have the ability to give other Snowflake users outside of their role access to this data.

How should these requirements be met?

- A. Grant ANALYST\_ROLE OWNERSHIP on the database, but make sure that ANALYST\_ROLE does not have the MANAGE GRANTS privilege on the account.
- B. Grant SYSADMIN OWNERSHIP of the database, but grant the create schema privilege on the database to the ANALYST\_ROLE.
- C. Make every schema in the database a MANAGED ACCESS schema, owned by SYSADMIN, and grant create privileges on each schema to the ANALYST\_ROLE for each type of object that needs to be created. Most Voted
- D. Grant ANALYST\_ROLE OWNERSHIP on the database, but grant the OWNERSHIP ON FUTURE [object type]s in database privilege to SYSADMIN.

**Correct Answer:** C

What considerations need to be taken when using database cloning as a tool for data lifecycle management in a development environment? (Choose two.)

- A. Any pipes in the source are not cloned.
- B. Any pipes in the source referring to internal stages are not cloned. Most Voted
- C. Any pipes in the source referring to external stages are not cloned.
- D. The clone inherits all granted privileges of all child objects in the source object, including the database.
- E. The clone inherits all granted privileges of all child objects in the source object, excluding the database. Most Voted

**Correct Answer:** BE

Which columns can be included in an external table schema? (Choose three.)

- A. VALUE Most Voted
- B. METADATA\$ROW\_ID
- C. METADATA\$ISUPDATE
- D. METADATA\$FILENAME Most Voted
- E. METADATA\$FILE\_ROW\_NUMBER Most Voted
- F. METADATA\$EXTERNAL\_TABLE\_PARTITION

**Correct Answer:** ADE

Which SQL ALTER command will MAXIMIZE memory and compute resources for a Snowpark stored procedure when executed on the snowpark\_opt\_wh warehouse?

- A. alter warehouse snowpark\_opt\_wh set max\_concurrency\_level = 1; Most Voted
- B. alter warehouse snowpark\_opt\_wh set max\_concurrency\_level = 2;
- C. alter warehouse snowpark\_opt\_wh set max\_concurrency\_level = 8;
- D. alter warehouse snowpark\_opt\_wh set max\_concurrency\_level = 16;

**Correct Answer:** A

An Architect clones a database and all of its objects, including tasks. After the cloning, the tasks stop running.

Why is this occurring?

- A. Tasks cannot be cloned.
- B. The objects that the tasks reference are not fully qualified.
- C. Cloned tasks are suspended by default and must be manually resumed. Most Voted
- D. The Architect has insufficient privileges to alter tasks on the cloned database.

**Correct Answer:** C

Which technique will efficiently ingest and consume semi-structured data for Snowflake data lake workloads?

- A. IDEF1X
- B. Schema-on-write
- C. Schema-on-read Most Voted
- D. Information schema

**Correct Answer:** C

Is it possible for a data provider account with a Snowflake Business Critical edition to share data with an Enterprise edition data consumer account?

- A. A Business Critical account cannot be a data sharing provider to an Enterprise consumer. Any consumer accounts must also be Business Critical.
- B. If a user in the provider account with role authority to CREATE or ALTER SHARE adds an Enterprise account as a consumer, it can import the share.
- C. If a user in the provider account with a share owning role sets SHARE\_RESTRICTIONS to False when adding an Enterprise consumer account, it can import the share.
- D. If a user in the provider account with a share owning role which also has OVERRIDE SHARE RESTRICTIONS privilege SHARE\_RESTRICTIONS set to False when adding an Enterprise consumer account, it can import the share. Most Voted

**Correct Answer:** D

Which of the following ingestion methods can be used to load near real-time data by using the messaging services provided by a cloud provider?

- A. Snowflake Connector for Kafka
- B. Snowflake streams
- C. Snowpipe Most Voted
- D. Spark

**Correct Answer:** C

An Architect is designing a file ingestion recovery solution. The project will use an internal named stage for file storage. Currently, in the case of an ingestion failure, the Operations team must manually download the failed file and check for errors.

Which downloading method should the Architect recommend that requires the LEAST amount of operational overhead?

- A. Use the Snowflake Connector for Python, connect to remote storage and download the file.
- B. Use the GET command in SnowSQL to retrieve the file.
- C. Use the GET command in Snowsight to retrieve the file.
- D. Use the Snowflake API endpoint and download the file.

**Correct Answer:** B

A table for IOT devices that measures water usage is created. The table quickly becomes large and contains more than 2 billion rows.

```
create table water_iot (
    UniqueId number,
    DeviceId varchar(20),
    DeviceManufacturer varchar(50)
    CustomerId varchar(20),
    IOT_timestamp timestamp_ntz,
    City varchar(80),
    Location varchar(50)
)
```

The general query patterns for the table are:

1. DeviceId, IOT\_timestamp and CustomerId are frequently used in the filter predicate for the select statement
2. The columns City and DeviceManufacturer are often retrieved
3. There is often a count on UniqueId

Which field(s) should be used for the clustering key?

- A. IOT\_timestamp
- B. City and DeviceManufacturer
- C. DeviceId and CustomerId Most Voted
- D. UniqueId

**Correct Answer: C**

Which Snowflake objects can be used in a data share? (Choose two.)

- A. Standard view
- B. Secure view
- C. Stored procedure
- D. External table
- E. Stream

**Correct Answer: BD**

A company has an external vendor who puts data into Google Cloud Storage. The company's Snowflake account is set up in Azure.

What would be the MOST efficient way to load data from the vendor into Snowflake?

- A. Ask the vendor to create a Snowflake account, load the data into Snowflake and create a data share.
- B. Create an external stage on Google Cloud Storage and use the external table to load the data into Snowflake. Most Voted
- C. Copy the data from Google Cloud Storage to Azure Blob storage using external tools and load data from Blob storage to Snowflake.
- D. Create a Snowflake Account in the Google Cloud Platform (GCP), ingest data into this account and use data replication to move the data from GCP to Azure.

**Correct Answer:** *B*

How can the Snowpipe REST API be used to keep a log of data load history?

- A. Call insertReport every 20 minutes, fetching the last 10,000 entries.
- B. Call loadHistoryScan every minute for the maximum time range.
- C. Call insertReport every 8 minutes for a 10-minute time range. Most Voted
- D. Call loadHistoryScan every 10 minutes for a 15-minute time range.

**Correct Answer:** *C*

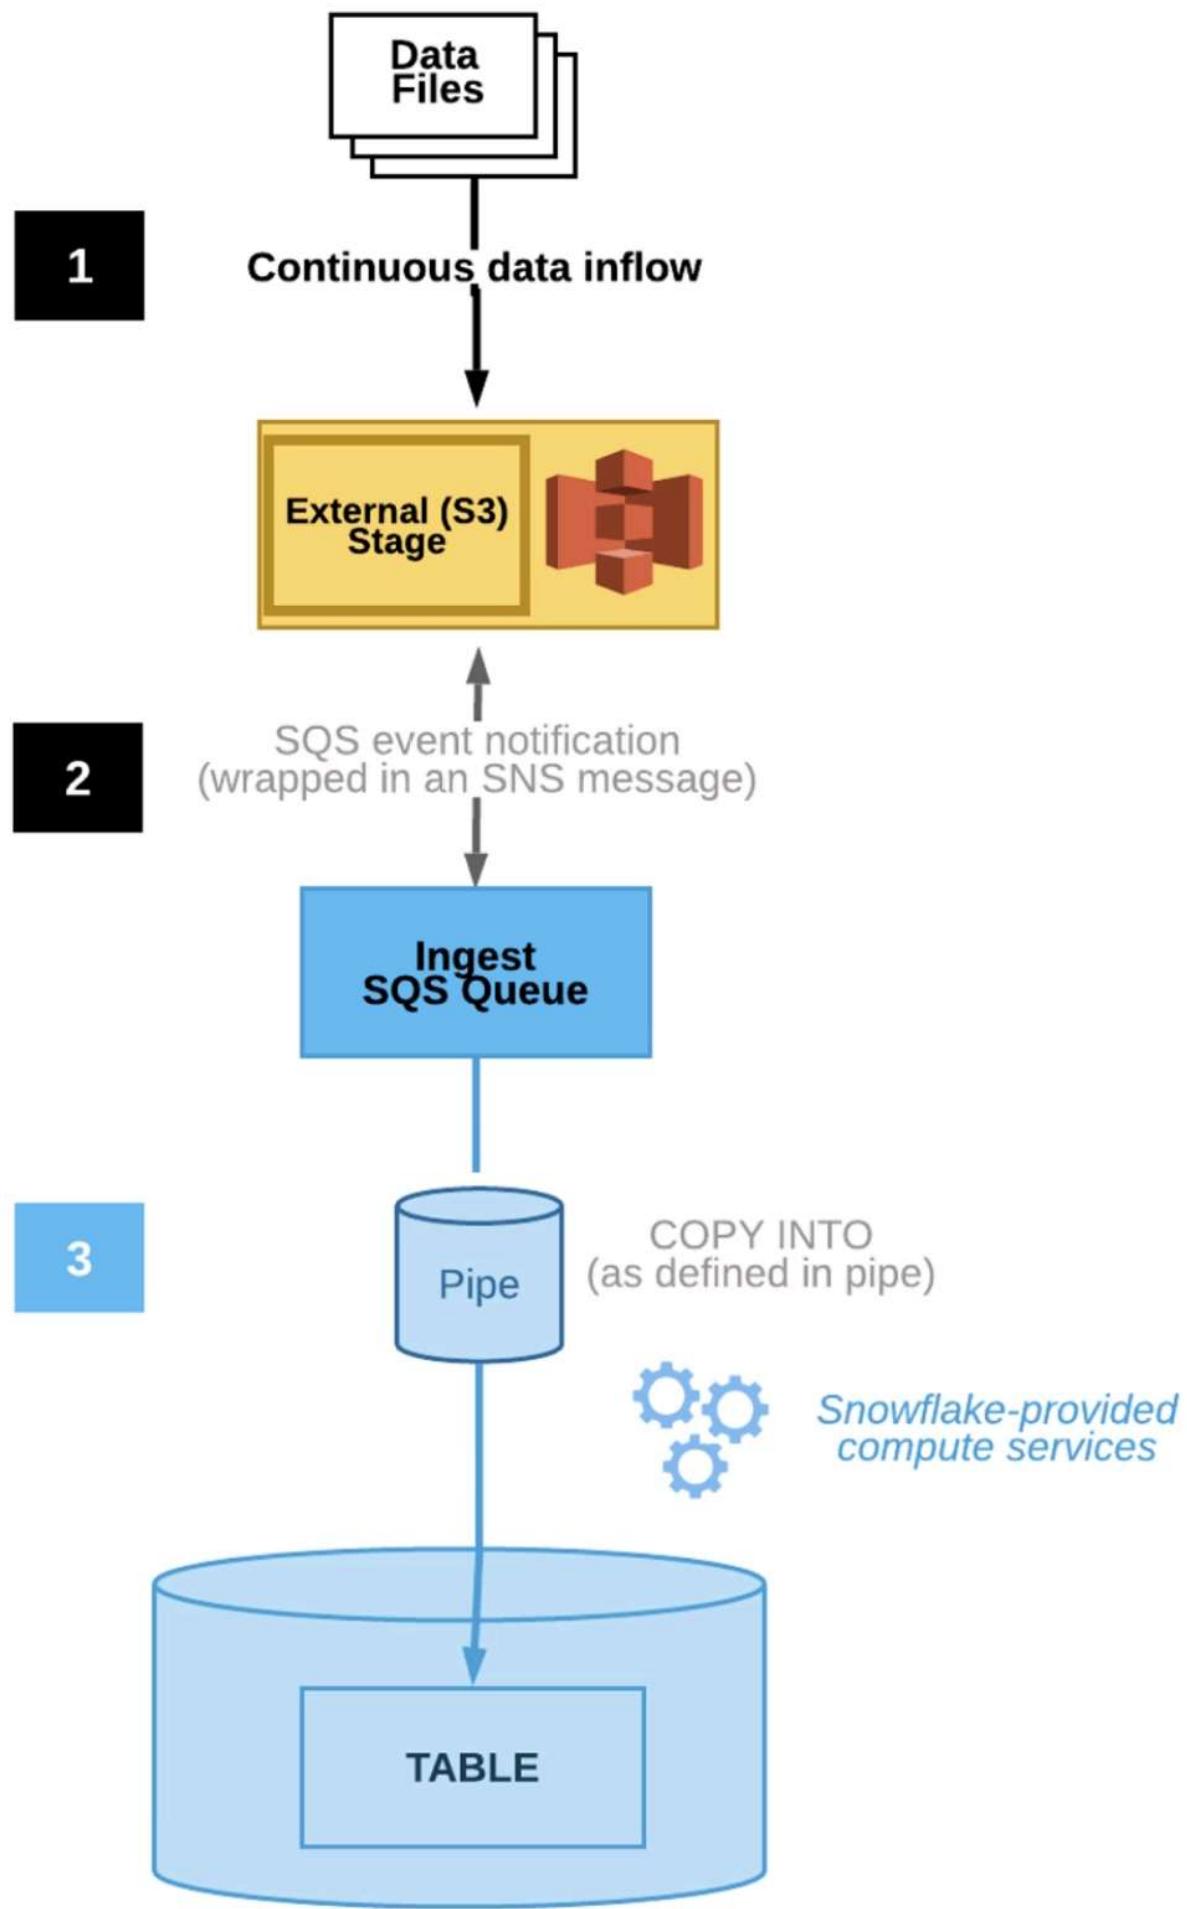
The diagram shows the process flow for Snowpipe auto-ingest with Amazon Simple Notification Service (SNS) with the following steps:

Step 1: Data files are loaded in a stage.

Step 2: An Amazon S3 event notification, published by SNS, informs Snowpipe – by way of Amazon Simple Queue Service (SQS) - that files are ready to load. Snowpipe copies the files into a queue.

Step 3: A Snowflake-provided virtual warehouse loads data from the queued files into the target table based on parameters defined in the specified pipe.

If an AWS Administrator accidentally deletes the SQS subscription to the SNS topic in Step 2, what will happen to the pipe that references the topic to receive event messages from Amazon



- A. The pipe will continue to receive the messages as Snowflake will automatically restore the subscription to the same SNS topic and will recreate the pipe by specifying the same SNS topic name in the pipe definition.
- B. The pipe will no longer be able to receive the messages and the user must wait for 24 hours from the time when the SNS topic subscription was deleted. Pipe recreation is not required as the pipe will reuse the same subscription to the existing SNS topic after 24 hours.
- C. The pipe will continue to receive the messages as Snowflake will automatically restore the subscription by creating a new SNS topic. Snowflake will then recreate the pipe by specifying the new SNS topic name in the pipe definition.
- D. The pipe will no longer be able to receive the messages. To restore the system immediately, the user needs to manually create a new SNS topic with a different name and then recreate the pipe by specifying the new SNS topic name in the pipe definition.

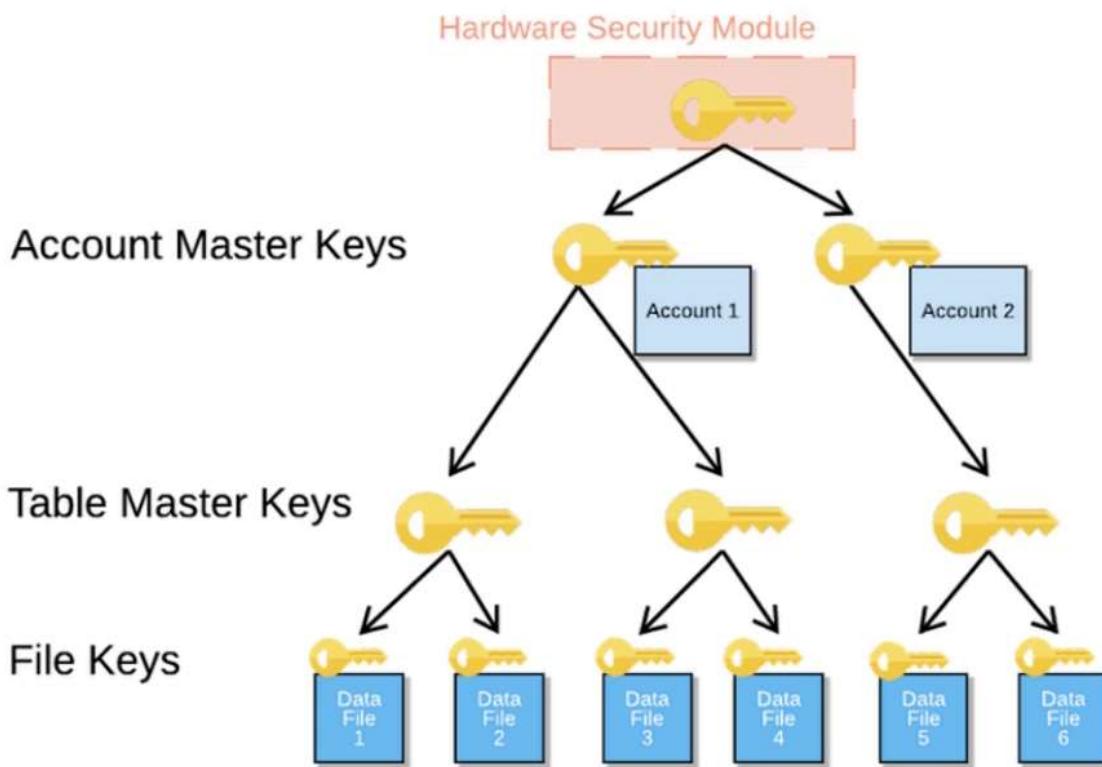
**Correct Answer:** D

An Architect needs to meet a company requirement to ingest files from the company's AWS storage accounts into the company's Snowflake account.

How can the ingestion of these files into the company's Snowflake account be initiated? (Choose two.)

- A. Configure the client application to call the Snowpipe REST endpoint when new files have arrived in Amazon S3 storage.
- B. Configure the client application to call the Snowpipe REST endpoint when new files have arrived in Amazon S3 Glacier storage.
- C. Create an AWS Lambda function to call the Snowpipe REST endpoint when new files have arrived in Amazon S3 storage.
- D. Configure AWS Simple Notification Service (SNS) to notify Snowpipe when new files have arrived in Amazon S3 storage.
- E. Configure the client application to issue a COPY INTO command to Snowflake when new files have arrived in Amazon S3 Glacier storage.

When activating Tri-Secret Secure in a hierarchical encryption model in a Snowflake account, at what level is the customer-managed key used?



- A. At the root level (HSM)
- B. At the account level (AMK)**
- C. At the table level (TMK)
- D. At the micro-partition level

**Correct Answer: B**

What are characteristics of the use of transactions in Snowflake? (Choose two.)

- A. Explicit transactions can contain DDL, DML, and query statements.
- B. The AUTOCOMMIT setting can be changed inside a stored procedure.
- C. A transaction can be started explicitly by executing a BEGIN WORK statement and end explicitly by executing a COMMIT WORK statement.
- D. A transaction can be started explicitly by executing a BEGIN TRANSACTION statement and end explicitly by executing an END TRANSACTION statement.
- E. Explicit transactions should contain only DML statements and query statements. All DDL statements implicitly commit active transactions.

**Correct Answer:** C,E

Which query will identify the specific days and virtual warehouses that would benefit from a multi-cluster warehouse to improve the performance of a particular workload?

- A. 

```
SELECT TO_DATE(START_TIME) AS DATE,
WAREHOUSE_NAME,
BYTES_SCANNED,
BYTES_SPILLED
FROM "SNOWFLAKE"."ACCOUNT_USAGE"."QUERY_HISTORY"
HAVING BYTES_SPILLED > BYTES_SCANNED;
```
- B. 

```
SELECT TO_DATE(START_TIME) AS DATE,
WAREHOUSE_NAME,
SUM(AVG_RUNNING) AS SUM_RUNNING,
SUM(AVG_QUEUED_LOAD) AS SUM_QUEUED
FROM "SNOWFLAKE"."ACCOUNT_USAGE"."WAREHOUSE_LOAD_HISTORY"
GROUP BY 1,2
HAVING SUM(AVG_QUEUED_LOAD) > 0;
```
- C. 

```
SELECT TO_DATE(START_TIME) AS DATE,
WAREHOUSE_NAME,
BYTES_SCANNED,
BYTES_SPILLED
FROM "SNOWFLAKE"."ACCOUNT_USAGE"."WAREHOUSE_LOAD_HISTORY"
HAVING BYTES_SPILLED > BYTES_SCANNED;
```
- D. 

```
SELECT TO_DATE(START_TIME) AS DATE,
WAREHOUSE_NAME,
BYTES_SPILLED_TO_LOCAL_STORAGE,
SUM(AVG_QUEUED_LOAD) AS SUM_QUEUED
FROM "SNOWFLAKE"."ACCOUNT_USAGE"."QUERY_HISTORY"
HAVING BYTES_SPILLED_TO_LOCAL_STORAGE > 0;
```

**Correct Answer:** B

A company is designing high availability and disaster recovery plans and needs to maximize redundancy and minimize recovery time objectives for their critical application processes.

Cost is not a concern as long as the solution is the best available.

The plan so far consists of the following steps:

1. Deployment of Snowflake accounts on two different cloud providers.
2. Selection of cloud provider regions that are geographically far apart.
3. The Snowflake deployment will replicate the databases and account data between both cloud provider accounts.
4. Implementation of Snowflake client redirect.

What is the MOST cost-effective way to provide the HIGHEST uptime and LEAST application disruption if there is a service event?

A. Connect the applications using the - URL.

Use the Business Critical Snowflake edition.

B. Connect the applications using the - URL.

Use the Virtual Private Snowflake (VPS) edition.

C. Connect the applications using the - URL.

Use the Enterprise Snowflake edition.

D. Connect the applications using the - URL.

Use the Business Critical Snowflake edition.

**Correct Answer: A**

A company's Architect needs to find an efficient way to get data from an external partner, who is also a Snowflake user. The current solution is based on daily JSON extracts that are placed on an FTP server and uploaded to Snowflake manually. The files are changed several times each month, and the ingestion process needs to be adapted to accommodate these changes.

What would be the MOST efficient solution?

A. Ask the partner to create a share and add the company's account.

B. Ask the partner to use the data lake export feature and place the data into cloud storage where Snowflake can natively ingest it (schema-on-read).

C. Keep the current structure but request that the partner stop changing files, instead only appending new files.

D. Ask the partner to set up a Snowflake reader account and use that account to get the data for ingestion.

**Correct Answer: A**

Based on the architecture in the image, how can the data from DB1 be copied into TBL2? (Choose two.)

Databases > DB1

Stage	Schema	Location
STAGE1	SH1	Snowflake

Databases > DB2

File Format	Schema	Type
FF_PIPE_1	SH1	CSV

- ```
use database DB1;
use schema SH1;
```
- A. `copy into DB1.SH2.TBL2
 from @STAGE1
 file_format = (format_name = FF_PIPE_1);`
- ```
use database DB1;
use schema SH2;
```
- B. `copy into TBL2
 from @STAGE1
 file_format = (format_name = FF_PIPE_1);`
- ```
use database DB1;
use schema SH2;
```
- C. `copy into DB1.SH2.TBL2
 from @STAGE1
 file_format = (format_name = DB1.SH1.FF_PIPE_1);`
- ```
use database DB1;
use schema SH2;
```
- D. `copy into DB1.SH2.TBL2
 from @DB1.SH1.STAGE1
 file_format = (format_name = FF_PIPE_1);`
- ```
use database DB1;
use schema SH2;
```
- E. `copy into TBL2
 from @DB1.SH1.STAGE1
 file_format = (format_name = DB1.SH1.FF_PIPE_1);`

Correct Answer: AE

Why might a Snowflake Architect use a star schema model rather than a 3NF model when designing a data architecture to run in Snowflake? (Choose two.)

- A. Snowflake cannot handle the joins implied in a 3NF data model.
- B. The Architect wants to remove data duplication from the data stored in Snowflake.
- C. The Architect is designing a landing zone to receive raw data into Snowflake.
- D. The BI tool needs a data model that allows users to summarize facts across different dimensions, or to drill down from the summaries.
- E. The Architect wants to present a simple flattened single view of the data to a particular group of end users.

**Correct Answer:** DE

An Architect is troubleshooting a query with poor performance using the QUERY\_HISTORY function. The Architect observes that the COMPILED\_TIME is greater than the EXECUTION\_TIME.

What is the reason for this?

- A. The query is processing a very large dataset.
- B. The query has overly complex logic.
- C. The query is queued for execution.
- D. The query is reading from remote storage.

**Correct Answer:** B

A Snowflake Architect is designing a multiple-account design strategy.

This strategy will be MOST cost-effective with which scenarios? (Choose two.)

- A. The company wants to clone a production database that resides on AWS to a development database that resides on Azure.
- B. The company needs to share data between two databases, where one must support Payment Card Industry Data Security Standard (PCI DSS) compliance but the other one does not.
- C. The company needs to support different role-based access control features for the development, test, and production environments.
- D. The company security policy mandates the use of different Active Directory instances for the development, test, and production environments.
- E. The company must use a specific network policy for certain users to allow and block given IP addresses.

**Correct Answer:** BD

The following table exists in the production database:

Name: user\_events

| Column          | Description                          |
|-----------------|--------------------------------------|
| user_hash_key   | hashed key with no meaning           |
| username        | user name                            |
| event_type      | the type of event that occurred      |
| event_timestamp | the date and time the event occurred |

A regulatory requirement states that the company must mask the username for events that are older than six months based on the current date when the data is queried.

How can the requirement be met without duplicating the event data and making sure it is applied when creating views using the table or cloning the table?

- A. Use a masking policy on the username column using an entitlement table with valid dates.
- B. Use a row level policy on the user\_events table using an entitlement table with valid dates.
- C. Use a masking policy on the username column with event\_timestamp as a conditional column.
- D. Use a secure view on the user\_events table using a case statement on the username column.

**Correct Answer: C**

Which data models can be used when modeling tables in a Snowflake environment? (Choose three.)

- A. Graph model
- B. Dimensional/Kimball
- C. Data lake
- D. Inmon/3NF
- E. Bayesian hierarchical model
- F. Data vault

**Correct Answer: BDF**

A Snowflake Architect is setting up database replication to support a disaster recovery plan. The primary database has external tables.

How should the database be replicated?

- A. Create a clone of the primary database then replicate the database.
- B. Move the external tables to a database that is not replicated, then replicate the primary database.
- C. Replicate the database ensuring the replicated database is in the same region as the external tables.
- D. Share the primary database with an account in the same region that the database will be replicated to.

**Correct Answer: B**

An Architect is integrating an application that needs to read and write data to Snowflake without installing any additional software on the application server.

How can this requirement be met?

- A. Use SnowSQL.
- B. Use the Snowpipe REST API.
- C. Use the Snowflake SQL REST API.
- D. Use the Snowflake ODBC driver.

**Correct Answer: C**

What transformations are supported in the below SQL statement? (Choose three.)

CREATE PIPE ... AS COPY ... FROM (...)

- A. Data can be filtered by an optional WHERE clause.
- B. Columns can be reordered.
- C. Columns can be omitted.
- D. Type casts are supported.
- E. Incoming data can be joined with other tables.
- F. The ON\_ERROR - ABORT\_STATEMENT command can be used.

**Correct Answer: BCD**

Data is being imported and stored as JSON in a VARIANT column. Query performance was fine, but most recently, poor query performance has been reported.

What could be causing this?

- A. There were JSON nulls in the recent data imports.
- B. The order of the keys in the JSON was changed.
- C. The recent data imports contained fewer fields than usual.
- D. There were variations in string lengths for the JSON values in the recent data imports.

**Correct Answer: A**

What step will improve the performance of queries executed against an external table?

- A. Partition the external table.
- B. Shorten the names of the source files.
- C. Convert the source files' character encoding to UTF-8.
- D. Use an internal stage instead of an external stage to store the source files.

**Correct Answer: A**

The Business Intelligence team reports that when some team members run queries for their dashboards in parallel with others, the query response time is getting significantly slower.

What can a Snowflake Architect do to identify what is occurring and troubleshoot this issue?

- A. Use larger warehouses to speed up the queries running in parallel. Identify the queries running in parallel using this query:

```
SELECT QUERY_ID
,USER_NAME
,WAREHOUSE_NAME
,WAREHOUSE_SIZE
,BYTES_SCANNED
,BYTES_SPILLED_TO_REMOTE_STORAGE
,BYTES_SPILLED_TO_REMOTE_STORAGE / BYTES_SCANNED AS SPILLING_READ_RATIO
FROM "SNOWFLAKE"."ACCOUNT_USAGE"."QUERY_HISTORY"
WHERE BYTES_SPILLED_TO_REMOTE_STORAGE > BYTES_SCANNed * 5
ORDER BY SPILLING_READ_RATIO DESC ;
```

- B. Increase the size of the warehouse cache to speed up concurrent queries. Identify the concurrent queries using this query:

```
SELECT WAREHOUSE_NAME ,COUNT(*) AS QUERY_COUNT ,SUM(BYTES_SCANNED) AS BYTES_SCANNED
,SUM(BYTES_SCANNED*PERCENTAGE_SCANNED_FROM_CACHE) AS BYTES_SCANNED_FROM_CACHE ,SUM(BYTES_SCANNED*PERCENTAGE_SCANNED_FROM_CACHE)
/ SUM(BYTES_SCANNED) AS PERCENT_SCANNED_FROM_CACHE
FROM "SNOWFLAKE"."ACCOUNT_USAGE"."QUERY_HISTORY"
WHERE START_TIME >= DATEADD(month,-1,current_timestamp()) AND BYTES_SCANNED > 0 GROUP BY 1 ORDER BY 5 ;
```

- C. Introduce multi-cluster warehouses to help with concurrent queries. Identify the concurrent queries by running this query:

```
SELECT TO_DATE(START_TIME) AS DATE
,WAREHOUSE_NAME ,SUM(AVG_RUNNING) AS SUM_RUNNING ,SUM(AVG_QUEUED_LOAD) AS SUM_QUEUED
FROM "SNOWFLAKE"."ACCOUNT_USAGE"."WAREHOUSE_LOAD_HISTORY" WHERE TO_DATE(START_TIME) >=
DATEADD(month,-1,CURRENT_TIMESTAMP()) GROUP BY 1,2 HAVING SUM(AVG_QUEUED_LOAD) >0 ;
```

- D. Identify which queries are spilled to remote storage and change the warehouse parameters to address this issue. Identify the issue by running this query:

```
SELECT QUERY_ID
,SUBSTR(QUERY_TEXT, 1, 50) PARTIAL_QUERY_TEXT
,USER_NAME
,WAREHOUSE_NAME
,WAREHOUSE_SIZE
,BYTES_SPILLED_TO_REMOTE_STORAGE
,START_TIME, END_TIME
,TOTAL_ELAPSED_TIME/1000 TOTAL_ELAPSED_TIME
FROM SNOWFLAKE.ACCOUNT_USAGE.QUERY_HISTORY WHERE BYTES_SPILLED_TO_REMOTE_STORAGE > 0 AND START_TIME::DATE > DATEADD('DAYS', -45, CURRENT_DATE) ORDER BY BYTES_SPILLED_TO_REMOTE_STORAGE DESC LIMIT 10 ;
```

**Correct Answer: C**

What is a key consideration when setting up search optimization service for a table?

- A. Search optimization service works best with a column that has a minimum of 100 K distinct values.
- B. Search optimization service can significantly improve query performance on partitioned external tables.
- C. Search optimization service can help to optimize storage usage by compressing the data into a GZIP format.
- D. The table must be clustered with a key having multiple columns for effective search optimization.

**Correct Answer: B**

A retail company has 2000+ stores spread across the country. Store Managers report that they are having trouble running key reports related to inventory management, sales targets, payroll, and staffing during business hours. The Managers report that performance is poor and time-outs occur frequently.

Currently all reports share the same Snowflake virtual warehouse.

How should this situation be addressed? (Choose two.)

- A. Use a Business Intelligence tool for in-memory computation to improve performance.
- B. Configure a dedicated virtual warehouse for the Store Manager team.
- C. Configure the virtual warehouse to be multi-clustered.
- D. Configure the virtual warehouse to size 4-XL.
- E. Advise the Store Manager team to defer report execution to off-business hours.

**Correct Answer:** BC

A company needs to have the following features available in its Snowflake account:

1. Support for Multi-Factor Authentication (MFA)
2. A minimum of 2 months of Time Travel availability
3. Database replication in between different regions
4. Native support for JDBC and ODBC
5. Customer-managed encryption keys using Tri-Secret Secure
6. Support for Payment Card Industry Data Security Standards (PCI DSS)

In order to provide all the listed services, what is the MINIMUM Snowflake edition that should be selected during account creation?

- A. Standard
- B. Enterprise
- C. Business Critical
- D. Virtual Private Snowflake (VPS)

**Correct Answer:** C

When using the COPY INTO [table] command with the CSV file format, how does the MATCH\_BY\_COLUMN\_NAME parameter behave?

- A. It expects a header to be present in the CSV file, which is matched to a case-sensitive table column name.
- B. The parameter will be ignored.
- C. The command will return an error.
- D. The command will return a warning stating that the file has unmatched columns.

**Correct Answer:** A

How can the Snowflake context functions be used to help determine whether a user is authorized to see data that has column-level security enforced? (Choose two.)

- A. Set masking policy conditions using CURRENT\_ROLE targeting the role in use for the current session.
- B. Set masking policy conditions using IS\_ROLE\_IN\_SESSION targeting the role in use for the current account.
- C. Set masking policy conditions using INVOKER\_ROLE targeting the executing role in a SQL statement.
- D. Determine if there are OWNERSHIP privileges on the masking policy that would allow the use of any function.
- E. Assign the ACCOUNTADMIN role to the user who is executing the object.

**Correct Answer:** AC

A retail company has over 3000 stores all using the same Point Of Sale (POS) system. The company wants to deliver near real-time sales results to category managers. The stores operate in a variety of time zones and exhibit a dynamic range of transactions each minute, with some stores having higher sales volumes than others.

Sales results are provided in a uniform fashion using data engineered fields that will be calculated in a complex data pipeline. Calculations include exceptions, aggregations, and scoring using external functions interfaced to scoring algorithms. The source data for aggregations has over 100M rows.

Every minute, the POS sends all sales transactions files to a cloud storage location with a naming convention that includes store numbers and timestamps to identify the set of transactions contained in the files. The files are typically less than 10MB in size.

How can the near real-time results be provided to the category managers? (Choose two.)

- A. All files should be concatenated before ingestion into Snowflake to avoid micro-ingestion.
- B. A Snowpipe should be created and configured with AUTO\_INGEST = TRUE. A stream should be created to process INSERTS into a single target table using the stream metadata to inform the store number and timestamps.
- C. A STREAM should be created to accumulate the near real-time data and a TASK should be created that runs at a frequency that matches the real-time analytics needs.
- D. An external scheduler should examine the contents of the cloud storage location and issue SnowSQL commands to process the data at a frequency that matches the real-time analytics needs.
- E. The COPY INTO command with a task scheduled to run every second should be used to achieve the near-real time requirement.

**Correct Answer:** BC

◀ Previous Questions

Next Questions ➔

Browse atleast 50% to increase passing rate :💡



Viewing page 2 out of 3 pages.

Viewing questions 51-100 out of 109 questions